

```

1 import numpy as np
2
3 from .layers import *
4 from .layer_utils import *
5
6 def softmax(y):
7     return np.exp(y)/np.reshape(np.sum(np.exp(y), axis=1), (len(y), 1))
8
9 class TwoLayerNet(object):
10     """
11     A two-layer fully-connected neural network with ReLU nonlinearity and
12     softmax loss that uses a modular layer design. We assume an input dimension
13     of D, a hidden dimension of H, and perform classification over C classes.
14
15     The architecture should be affine - relu - affine - softmax.
16
17     Note that this class does not implement gradient descent; instead, it
18     will interact with a separate Solver object that is responsible for running
19     optimization.
20
21     The learnable parameters of the model are stored in the dictionary
22     self.params that maps parameter names to numpy arrays.
23     """
24
25     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
26                 dropout=0, weight_scale=1e-3, reg=0.0):
27         """
28         Initialize a new network.
29
30         Inputs:
31         - input_dim: An integer giving the size of the input
32         - hidden_dims: An integer giving the size of the hidden layer
33         - num_classes: An integer giving the number of classes to classify
34         - dropout: Scalar between 0 and 1 giving dropout strength.
35         - weight_scale: Scalar giving the standard deviation for random
36           initialization of the weights.
37         - reg: Scalar giving L2 regularization strength.
38         """
39         self.params = {}
40         self.reg = reg
41
42         # ===== #
43         # YOUR CODE HERE:
44         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
45         # self.params['W2'], self.params['b1'] and self.params['b2']. The
46         # biases are initialized to zero and the weights are initialized
47         # so that each parameter has mean 0 and standard deviation weight_scale.
48         # The dimensions of W1 should be (input_dim, hidden_dim) and the
49         # dimensions of W2 should be (hidden_dims, num_classes)
50         # ===== #
51
52         self.params['W1'] = weight_scale * np.random.randn(input_dim, hidden_dims)
53         self.params['b1'] = np.zeros(hidden_dims)
54         self.params['W2'] = weight_scale * np.random.randn(hidden_dims, num_classes)
55         self.params['b2'] = np.zeros(num_classes)
56
57         # ===== #
58         # END YOUR CODE HERE
59         # ===== #

```

```

60
61 def loss(self, X, y=None):
62     """
63     Compute loss and gradient for a minibatch of data.
64
65     Inputs:
66     - X: Array of input data of shape (N, d_1, ..., d_k)
67     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
68
69     Returns:
70     If y is None, then run a test-time forward pass of the model and return:
71     - scores: Array of shape (N, C) giving classification scores, where
72       scores[i, c] is the classification score for X[i] and class c.
73
74     If y is not None, then run a training-time forward and backward pass and
75     return a tuple of:
76     - loss: Scalar value giving the loss
77     - grads: Dictionary with the same keys as self.params, mapping parameter
78       names to gradients of the loss with respect to those parameters.
79     """
80     scores = None
81
82     # ===== #
83     # YOUR CODE HERE:
84     # Implement the forward pass of the two-layer neural network. Store
85     # the class scores as the variable 'scores'. Be sure to use the layers
86     # you prior implemented.
87     # ===== #
88     h1, cache1 = affine_relu_forward(X, self.params['W1'], self.params['b1'])
89     scores, cache2 = affine_forward(h1, self.params['W2'], self.params['b2'])
90     # ===== #
91     # END YOUR CODE HERE
92     # ===== #
93
94     # If y is None then we are in test mode so just return scores
95     if y is None:
96         return scores
97
98     loss, grads = 0, {}
99     # ===== #
100    # YOUR CODE HERE:
101    # Implement the backward pass of the two-layer neural net. Store
102    # the loss as the variable 'loss' and store the gradients in the
103    # 'grads' dictionary. For the grads dictionary, grads['W1'] holds
104    # the gradient for W1, grads['b1'] holds the gradient for b1, etc.
105    # i.e., grads[k] holds the gradient for self.params[k].
106    #
107    # Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
108    # for each W. Be sure to include the 0.5 multiplying factor to
109    # match our implementation.
110    #
111    # And be sure to use the layers you prior implemented.
112    # ===== #
113    loss, dout = softmax_loss(scores, y)
114    loss += 0.5*self.reg * (np.linalg.norm(self.params['W1'])**2 +
np.linalg.norm(self.params['W2'])**2)
115    dout, grads['W2'], grads['b2'] = affine_backward(dout, cache2)
116    grads['W2'] += self.reg * self.params['W2']
117    dout, grads['W1'], grads['b1'] = affine_relu_backward(dout, cache1)
118    grads['W1'] += self.reg * self.params['W1']

```

```

119
120 # ===== #
121 # END YOUR CODE HERE
122 # ===== #
123
124     return loss, grads
125
126
127 class FullyConnectedNet(object):
128     """
129     A fully-connected neural network with an arbitrary number of hidden layers,
130     ReLU nonlinearities, and a softmax loss function. This will also implement
131     dropout and batch normalization as options. For a network with L layers,
132     the architecture will be
133
134     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
135
136     where batch normalization and dropout are optional, and the {...} block is
137     repeated L - 1 times.
138
139     Similar to the TwoLayerNet above, learnable parameters are stored in the
140     self.params dictionary and will be learned using the Solver class.
141     """
142
143     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
144                 dropout=0, use_batchnorm=False, reg=0.0,
145                 weight_scale=1e-2, dtype=np.float32, seed=None):
146         """
147         Initialize a new FullyConnectedNet.
148
149         Inputs:
150         - hidden_dims: A list of integers giving the size of each hidden layer.
151         - input_dim: An integer giving the size of the input.
152         - num_classes: An integer giving the number of classes to classify.
153         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
154           the network should not use dropout at all.
155         - use_batchnorm: Whether or not the network should use batch normalization.
156         - reg: Scalar giving L2 regularization strength.
157         - weight_scale: Scalar giving the standard deviation for random
158           initialization of the weights.
159         - dtype: A numpy datatype object; all computations will be performed using
160           this datatype. float32 is faster but less accurate, so you should use
161           float64 for numeric gradient checking.
162         - seed: If not None, then pass this random seed to the dropout layers. This
163           will make the dropout layers deterministic so we can gradient check the
164           model.
165         """
166         self.use_batchnorm = use_batchnorm
167         self.use_dropout = dropout > 0
168         self.reg = reg
169         self.num_layers = 1 + len(hidden_dims)
170         self.dtype = dtype
171         self.params = {}
172
173         # ===== #
174         # YOUR CODE HERE:
175         #   Initialize all parameters of the network in the self.params dictionary.
176         #   The weights and biases of layer 1 are W1 and b1; and in general the
177         #   weights and biases of layer i are Wi and bi. The
178         #   biases are initialized to zero and the weights are initialized

```

```

179 # so that each parameter has mean 0 and standard deviation weight_scale.
180 # ===== #
181
182 self.params['W1'] = weight_scale * np.random.randn(input_dim, hidden_dims[0])
183 self.params['b1'] = np.zeros(hidden_dims[0])
184 for i in np.arange(self.num_layers - 2):
185     self.params['W' + str(i+2)] = weight_scale * np.random.randn(hidden_dims[i],
hidden_dims[i+1])
186     self.params['b' + str(i+2)] = np.zeros(hidden_dims[i+1])
187 self.params['W' + str(self.num_layers)] = weight_scale *
np.random.randn(hidden_dims[-1], num_classes)
188 self.params['b' + str(self.num_layers)] = np.zeros(num_classes)
189
190 # ===== #
191 # END YOUR CODE HERE
192 # ===== #
193
194 # When using dropout we need to pass a dropout_param dictionary to each
195 # dropout layer so that the layer knows the dropout probability and the mode
196 # (train / test). You can pass the same dropout_param to each dropout layer.
197 self.dropout_param = {}
198 if self.use_dropout:
199     self.dropout_param = {'mode': 'train', 'p': dropout}
200     if seed is not None:
201         self.dropout_param['seed'] = seed
202
203 # With batch normalization we need to keep track of running means and
204 # variances, so we need to pass a special bn_param object to each batch
205 # normalization layer. You should pass self.bn_params[0] to the forward pass
206 # of the first batch normalization layer, self.bn_params[1] to the forward
207 # pass of the second batch normalization layer, etc.
208 self.bn_params = []
209 if self.use_batchnorm:
210     self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]
211
212 # Cast all parameters to the correct datatype
213 for k, v in self.params.items():
214     self.params[k] = v.astype(dtype)
215
216
217 def loss(self, X, y=None):
218     """
219     Compute loss and gradient for the fully-connected net.
220
221     Input / output: Same as TwoLayerNet above.
222     """
223     X = X.astype(self.dtype)
224     mode = 'test' if y is None else 'train'
225
226     # Set train/test mode for batchnorm params and dropout param since they
227     # behave differently during training and testing.
228     if self.dropout_param is not None:
229         self.dropout_param['mode'] = mode
230     if self.use_batchnorm:
231         for bn_param in self.bn_params:
232             bn_param['mode'] = mode
233
234     scores = None
235
236 # ===== #

```

```

237     # YOUR CODE HERE:
238     #     Implement the forward pass of the FC net and store the output
239     #     scores as the variable "scores".
240     # ===== #
241     h = X
242     caches = []
243     for i in np.arange(self.num_layers - 1):
244         h, cache = affine_relu_forward(h, self.params['W' + str(i+1)], self.params['b'
+ str(i+1)])
245         caches.append(cache)
246
247     scores, cache = affine_forward(h, self.params['W' + str(self.num_layers)],
self.params['b' + str(self.num_layers)])
248     caches.append(cache)
249     # ===== #
250     # END YOUR CODE HERE
251     # ===== #
252
253     # If test mode return early
254     if mode == 'test':
255         return scores
256
257     loss, grads = 0.0, {}
258     # ===== #
259     # YOUR CODE HERE:
260     #     Implement the backwards pass of the FC net and store the gradients
261     #     in the grads dict, so that grads[k] is the gradient of self.params[k]
262     #     Be sure your L2 regularization includes a 0.5 factor.
263     # ===== #
264
265     loss, dout = softmax_loss(scores, y)
266     dout, grads['W' + str(self.num_layers)], grads['b' + str(self.num_layers)] =
affine_backward(dout, caches[self.num_layers-1])
267     w_loss = np.linalg.norm(self.params['W' + str(self.num_layers)])**2
268     grads['W' + str(self.num_layers)] += self.reg * self.params['W' +
str(self.num_layers)]
269     for i in np.arange(self.num_layers-1):
270         w_loss += np.linalg.norm(self.params['W' + str(self.num_layers-i-1)])**2
271         dout, grads['W' + str(self.num_layers-i-1)], grads['b' + str(self.num_layers-
i-1)] = affine_relu_backward(dout, caches[self.num_layers-i-2])
272         grads['W' + str(self.num_layers-i-1)] += self.reg * self.params['W' +
str(self.num_layers-i-1)]
273         loss += 0.5*self.reg * w_loss
274
275
276     # ===== #
277     # END YOUR CODE HERE
278     # ===== #
279     return loss, grads
280

```