

```

1 import numpy as np
2 import pdb
3
4
5
6
7 def affine_forward(x, w, b):
8     """
9     Computes the forward pass for an affine (fully-connected) layer.
10
11     The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
12     examples, where each example x[i] has shape (d_1, ..., d_k). We will
13     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
14     then transform it to an output vector of dimension M.
15
16     Inputs:
17     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
18     - w: A numpy array of weights, of shape (D, M)
19     - b: A numpy array of biases, of shape (M,)
20
21     Returns a tuple of:
22     - out: output, of shape (N, M)
23     - cache: (x, w, b)
24     """
25
26     # ===== #
27     # YOUR CODE HERE:
28     #   Calculate the output of the forward pass. Notice the dimensions
29     #   of w are D x M, which is the transpose of what we did in earlier
30     #   assignments.
31     # ===== #
32
33     X = np.reshape(x, (x.shape[0], -1))
34     out = X@w + b
35
36     # ===== #
37     # END YOUR CODE HERE
38     # ===== #
39
40     cache = (x, w, b)
41     return out, cache
42
43
44 def affine_backward(dout, cache):
45     """
46     Computes the backward pass for an affine layer.
47
48     Inputs:
49     - dout: Upstream derivative, of shape (N, M)
50     - cache: Tuple of:
51       - x: Input data, of shape (N, d_1, ..., d_k)
52       - w: Weights, of shape (D, M)
53
54     Returns a tuple of:
55     - dx: Gradient with respect to x, of shape (N, d_1, ..., d_k)
56     - dw: Gradient with respect to w, of shape (D, M)
57     - db: Gradient with respect to b, of shape (M,)
58     """
59     x, w, b = cache

```

```

60 dx, dw, db = None, None, None
61
62 # ===== #
63 # YOUR CODE HERE:
64 #   Calculate the gradients for the backward pass.
65 # ===== #
66
67 # dout is N x M
68 # dx should be N x d1 x ... x dk; it relates to dout through multiplication with
w, which is D x M
69 # dw should be D x M; it relates to dout through multiplication with x, which is N
x D after reshaping
70 # db should be M; it is just the sum over dout examples
71
72 N = x.shape[0]
73 X = np.reshape(x, (N, -1))
74 dx = np.reshape(dout@w.T, x.shape)
75 dw = X.T@dout
76 db = (dout.T@np.ones((N,)))
77
78 # ===== #
79 # END YOUR CODE HERE
80 # ===== #
81
82 return dx, dw, db
83
84 def relu_forward(x):
85     """
86     Computes the forward pass for a layer of rectified linear units (ReLU).
87
88     Input:
89     - x: Inputs, of any shape
90
91     Returns a tuple of:
92     - out: Output, of the same shape as x
93     - cache: x
94     """
95     # ===== #
96     # YOUR CODE HERE:
97     #   Implement the ReLU forward pass.
98     # ===== #
99     out = np.maximum(x, 0)
100
101     # ===== #
102     # END YOUR CODE HERE
103     # ===== #
104
105     cache = x
106     return out, cache
107
108
109 def relu_backward(dout, cache):
110     """
111     Computes the backward pass for a layer of rectified linear units (ReLU).
112
113     Input:
114     - dout: Upstream derivatives, of any shape
115     - cache: Input x, of same shape as dout
116
117     Returns:

```

```

118 - dx: Gradient with respect to x
119 """
120 x = cache
121
122 # ===== #
123 # YOUR CODE HERE:
124 #   Implement the ReLU backward pass
125 # ===== #
126
127 # ReLU directs linearly to those > 0
128 dx = dout * (x>0)
129
130 # ===== #
131 # END YOUR CODE HERE
132 # ===== #
133
134 return dx
135
136
137 def softmax_loss(x, y):
138     """
139     Computes the loss and gradient for softmax classification.
140
141     Inputs:
142     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
143         for the ith input.
144     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
145         0 <= y[i] < C
146
147     Returns a tuple of:
148     - loss: Scalar giving the loss
149     - dx: Gradient of the loss with respect to x
150     """
151
152     probs = np.exp(x - np.max(x, axis=1, keepdims=True))
153     probs /= np.sum(probs, axis=1, keepdims=True)
154     N = x.shape[0]
155     loss = -np.sum(np.log(probs[np.arange(N), y])) / N
156     dx = probs.copy()
157     dx[np.arange(N), y] -= 1
158     dx /= N
159     return loss, dx
160

```