

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 def relu(x):
5     return np.maximum(0, x)
6 def softmax(y):
7     return np.exp(y)/np.reshape(np.sum(np.exp(y), axis=1), (len(y), 1))
8
9 class TwoLayerNet(object):
10     """
11     A two-layer fully-connected neural network. The net has an input dimension of
12     D, a hidden layer dimension of H, and performs classification over C classes.
13     We train the network with a softmax loss function and L2 regularization on the
14     weight matrices. The network uses a ReLU nonlinearity after the first fully
15     connected layer.
16
17     In other words, the network has the following architecture:
18
19     input - fully connected layer - ReLU - fully connected layer - softmax
20
21     The outputs of the second fully-connected layer are the scores for each class.
22     """
23
24     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
25         """
26         Initialize the model. Weights are initialized to small random values and
27         biases are initialized to zero. Weights and biases are stored in the
28         variable self.params, which is a dictionary with the following keys:
29
30         W1: First layer weights; has shape (H, D)
31         b1: First layer biases; has shape (H,)
32         W2: Second layer weights; has shape (C, H)
33         b2: Second layer biases; has shape (C,)
34
35         Inputs:
36         - input_size: The dimension D of the input data.
37         - hidden_size: The number of neurons H in the hidden layer.
38         - output_size: The number of classes C.
39         """
40         self.params = {}
41         self.params['W1'] = std * np.random.randn(hidden_size, input_size)
42         self.params['b1'] = np.zeros(hidden_size)
43         self.params['W2'] = std * np.random.randn(output_size, hidden_size)
44         self.params['b2'] = np.zeros(output_size)
45
46     def loss(self, X, y=None, reg=0.0):
47         """
48         Compute the loss and gradients for a two layer fully connected neural
49         network.
50
51         Inputs:
52         - X: Input data of shape (N, D). Each X[i] is a training sample.
53         - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
54             an integer in the range 0 <= y[i] < C. This parameter is optional; if it
55             is not passed then we only return scores, and if it is passed then we
56             instead return the loss and gradients.
57         - reg: Regularization strength.
58
59         Returns:

```

```

60 If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
61 the score for class c on input X[i].
62
63 If y is not None, instead return a tuple of:
64 - loss: Loss (data loss and regularization loss) for this batch of training
65 samples.
66 - grads: Dictionary mapping parameter names to gradients of those parameters
67 with respect to the loss function; has the same keys as self.params.
68 """
69 # Unpack variables from the params dictionary
70 W1, b1 = self.params['W1'], self.params['b1']
71 W2, b2 = self.params['W2'], self.params['b2']
72 N, D = X.shape
73
74 # Compute the forward pass
75 scores = None
76
77 # ===== #
78 # YOUR CODE HERE:
79 # Calculate the output scores of the neural network. The result
80 # should be (N, C). As stated in the description for this class,
81 # there should not be a ReLU layer after the second FC layer.
82 # The output of the second FC layer is the output scores. Do not
83 # use a for loop in your implementation.
84 # ===== #
85 v1 = X@W1.T + b1
86 h1 = relu(v1)
87 scores = h1@W2.T + b2
88
89 # ===== #
90 # END YOUR CODE HERE
91 # ===== #
92
93
94 # If the targets are not given then jump out, we're done
95 if y is None:
96     return scores
97
98 # Compute the loss
99 loss = None
100
101 # ===== #
102 # YOUR CODE HERE:
103 # Calculate the loss of the neural network. This includes the
104 # softmax loss and the L2 regularization for W1 and W2. Store the
105 # total loss in the variable loss. Multiply the regularization
106 # loss by 0.5 (in addition to the factor reg).
107 # ===== #
108
109 # scores is num_examples by num_classes
110 loss = 0.5*reg*(np.linalg.norm(W1)**2 + np.linalg.norm(W2)**2)
111 loss += np.sum(np.log(np.sum(np.exp(scores), axis = 1)) - scores[np.arange(N),
y])/N
112
113 # ===== #
114 # END YOUR CODE HERE
115 # ===== #
116
117 grads = {}
118

```

```

119 # ===== #
120 # YOUR CODE HERE:
121 # Implement the backward pass. Compute the derivatives of the
122 # weights and the biases. Store the results in the grads
123 # dictionary. e.g., grads['W1'] should store the gradient for
124 # W1, and be of the same size as W1.
125 # ===== #
126 dLds = softmax(scores)
127 dLds[np.arange(N), y] -= 1
128 dLds /= N
129 grads['W2'] = dLds.T@h1 + reg*W2
130 grads['b2'] = dLds.T@np.ones(len(h1))
131 dhdv = np.where(X@W1.T > 0, 1, 0)
132 dd = np.multiply(dhdv, dLds@W2)
133 grads['W1'] = dd.T @ X + reg*W1
134 grads['b1'] = dd.T @ np.ones(len(X))
135
136 # ===== #
137 # END YOUR CODE HERE
138 # ===== #
139
140 return loss, grads
141
142 def train(self, X, y, X_val, y_val,
143         learning_rate=1e-3, learning_rate_decay=0.95,
144         reg=1e-5, num_iters=100,
145         batch_size=200, verbose=False):
146     """
147     Train this neural network using stochastic gradient descent.
148
149     Inputs:
150     - X: A numpy array of shape (N, D) giving training data.
151     - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
152       X[i] has label c, where 0 ≤ c < C.
153     - X_val: A numpy array of shape (N_val, D) giving validation data.
154     - y_val: A numpy array of shape (N_val,) giving validation labels.
155     - learning_rate: Scalar giving learning rate for optimization.
156     - learning_rate_decay: Scalar giving factor used to decay the learning rate
157       after each epoch.
158     - reg: Scalar giving regularization strength.
159     - num_iters: Number of steps to take when optimizing.
160     - batch_size: Number of training examples to use per step.
161     - verbose: boolean; if true print progress during optimization.
162     """
163     num_train = X.shape[0]
164     iterations_per_epoch = max(num_train / batch_size, 1)
165
166     # Use SGD to optimize the parameters in self.model
167     loss_history = []
168     train_acc_history = []
169     val_acc_history = []
170
171     for it in np.arange(num_iters):
172         X_batch = None
173         y_batch = None
174
175         # ===== #
176         # YOUR CODE HERE:
177         # Create a minibatch by sampling batch_size samples randomly.
178         # ===== #

```

```

179     samps = np.random.choice(np.arange(len(X)), batch_size)
180     X_batch = X[samps]
181     y_batch = y[samps]
182     # ===== #
183     # END YOUR CODE HERE
184     # ===== #
185
186     # Compute loss and gradients using the current minibatch
187     loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
188     loss_history.append(loss)
189
190     # ===== #
191     # YOUR CODE HERE:
192     #     Perform a gradient descent step using the minibatch to update
193     #     all parameters (i.e., W1, W2, b1, and b2).
194     # ===== #
195
196     self.params['W1'] -= learning_rate * grads['W1']
197     self.params['b1'] -= learning_rate * grads['b1']
198     self.params['W2'] -= learning_rate * grads['W2']
199     self.params['b2'] -= learning_rate * grads['b2']
200
201     # ===== #
202     # END YOUR CODE HERE
203     # ===== #
204
205     if verbose and it % 100 == 0:
206         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
207
208     # Every epoch, check train and val accuracy and decay learning rate.
209     if it % iterations_per_epoch == 0:
210         # Check accuracy
211         train_acc = (self.predict(X_batch) == y_batch).mean()
212         val_acc = (self.predict(X_val) == y_val).mean()
213         train_acc_history.append(train_acc)
214         val_acc_history.append(val_acc)
215
216         # Decay learning rate
217         learning_rate *= learning_rate_decay
218
219     return {
220         'loss_history': loss_history,
221         'train_acc_history': train_acc_history,
222         'val_acc_history': val_acc_history,
223     }
224
225 def predict(self, X):
226     """
227     Use the trained weights of this two-layer network to predict labels for
228     data points. For each data point we predict scores for each of the C
229     classes, and assign each data point to the class with the highest score.
230
231     Inputs:
232     - X: A numpy array of shape (N, D) giving N D-dimensional data points to
233         classify.
234
235     Returns:
236     - y_pred: A numpy array of shape (N,) giving predicted labels for each of
237         the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
238         to have class c, where 0 <= c < C.

```

```

239     """
240     y_pred = None
241
242     # ===== #
243     # YOUR CODE HERE:
244     #   Predict the class given the input data.
245     # ===== #
246     W1 = self.params['W1']
247     b1 = self.params['b1']
248     W2 = self.params['W2']
249     b2 = self.params['b2']
250     v1 = X@W1.T + b1
251     h1 = relu(v1)
252     scores = h1@W2.T + b2
253     y_pred = np.argmax(softmax(scores), axis = 1)
254
255
256     # ===== #
257     # END YOUR CODE HERE
258     # ===== #
259
260     return y_pred
261

```